

COLD WAR COWBOYS

Content Authoring Manual

Scene Writing Reference for the NarrativeDirector Template System

Version 1.0 | May 2026
Prepared for Dan James | Robotics Division

A cyberpunk Crusader Kings / Telltale hybrid built in Blazor WebAssembly

Table of Contents

1. Overview: How CWC's Narrative System Works

Cold War Cowboys operates on three interlocking layers, each feeding the others. Understanding this architecture is essential before writing a single line of dialogue.

1.1 The Three Layers

Layer	Function
Strategy Sandbox	The player manages operatives, assigns missions, allocates resources, and navigates corporate politics. This layer generates the raw state: stats, relationships, corruption indices, mission outcomes.
Narrative Experience	Scenes trigger based on sandbox state. These are the moments the player reads, hears, and makes choices within. Dialogue, confrontation, introspection. This is where meaning is created.
Human Simulation	Operatives are not units. They accumulate stress, form bonds, lose loyalty, break. The narrative layer makes this visible; the simulation layer makes it real.

Your job as content author is the narrative layer. You write the scenes that make the sandbox state legible and emotionally resonant. The engine handles everything else.

1.2 How Scenes Trigger

The **NarrativeDirector** runs a check each game cycle. It evaluates every registered scene template against the current game state. When a template's trigger conditions are met, the scene enters a candidate pool. The Director selects from candidates based on priority, cooldowns, and narrative pacing rules.

The flow is deterministic from the author's perspective:

1. Trigger condition met (e.g., an operative hits `wet_work_count` ≥ 3 AND `conscience` ≥ 40)
2. Cast resolution: NarrativeDirector fills role slots with operatives matching the criteria
3. Scene renders in the UI with the resolved cast, their portraits, and mood tags
4. Player reads dialogue, makes a choice from 2-4 options
5. Effects applied: stat deltas, flag settings, corporate standing changes
6. Flags set to prevent re-triggering and to gate future scenes

You define steps 1, 3, 4, and 5. The engine handles 2 and 6 automatically.

1.3 Role-Based Casting

Scenes do not reference specific operatives by name. Instead, they reference **roles** that the NarrativeDirector fills dynamically based on operative fitness. This means one scene template can play out with entirely different casts across different playthroughs.

Role	Archetype	Fitness Criteria
conscience	The one who still feels it	Highest conscience stat among eligible operatives
mirror	The observer who names what's happening	High social + perception; the witness to moral descent

weapon	The one who does the dirty work	Highest combat or wet_work_count; low conscience
innocent	The one who hasn't been corrupted yet	Lowest corruption_exposure; fewest missions
survivor	The one who's been through it and endured	High stress + high loyalty; still functional
climber	The corporate ladder aspirant	Highest ambition; responsive to board_confidence
anchor	The emotional tether to humanity	Highest relationship bond with triggering_operative
wildcard	Unpredictable element	Extreme stat variance; high volatility score

Design Note

Roles are tendencies, not fixed assignments. The NarrativeDirector re-evaluates fitness each cycle. An operative who is the 'conscience' in Act 1 might become the 'weapon' by Act 3. This is the game working as intended.

2. Scene File Format (JSON Reference)

Scene files are JSON documents consumed directly by the NarrativeDirector. They live in **Data/Templates/** and are hot-loadable: save the file, trigger the condition in-game, and the scene plays immediately. No code changes, no recompilation.

2.1 Annotated Example

Below is a complete scene file with every field documented:

```
{
  "id": "corruption_third_wet_job",
  "priority": 80,
  "cooldown_cycles": 0,
  "trigger": {
    "type": "any_operative",
    "conditions": {
      "wet_work_count": { "gte": 3 },
      "conscience": { "gte": 40 }
    }
  },
  "cast": {
    "protagonist": "triggering_operative",
    "observer": { "role": "mirror", "fallback": "highest_social" },
    "challenger": { "role": "conscience", "fallback": "highest_conscience" }
  },
  "nodes": [
    {
      "id": "opening",
      "speaker": "observer",
      "portrait_mood": "concerned",
      "text": "Three names on your list now. {gender:m|Man|Woman}...",
      "next": "choice_node"
    },
    {
      "id": "choice_node",
      "speaker": "protagonist",
      "portrait_mood": "neutral",
      "text": "The third was necessary. They all were.",
      "choices": [
        {
          "label": "They were all necessary.",

```

```

    "effects": [
      { "target": "protagonist", "stat": "conscience", "delta": -8 },
      { "target": "protagonist", "stat": "stress", "delta": -3 }
    ],
    "set_flags": ["accepted_violence"],
    "next": "cold_resolution"
  },
  {
    "label": "I need to talk to someone.",
    "effects": [
      { "target": "protagonist", "stat": "stress", "delta": -10 },
      { "target": "protagonist", "stat": "loyalty", "delta": 3 }
    ],
    "set_flags": ["sought_help"],
    "corporate_standing": { "delta": -2 },
    "next": "humane_resolution"
  }
]
}
],
"on_complete_flags": ["third_wet_job_scene_played"]
}

```

2.2 Field Reference

Field	Type	Description
id	string	Unique scene identifier. Use snake_case. Must be globally unique across all templates.
priority	int (0-100)	Higher priority scenes are selected first when multiple triggers match in the same cycle.
cooldown_cycles	int	Minimum cycles before this scene can trigger again. Set 0 for one-shot scenes.
trigger.type	string	Who is checked: any_operative, specific_operative, relationship, corporate_state, post_mission.
trigger.conditions	object	Stat thresholds using gte, lte, eq, between operators. All conditions must be met (AND logic).
cast	object	Role-to-operative mapping. Keys are role names used in nodes; values are resolution rules.
nodes	array	Ordered dialogue nodes. Each has speaker, text, optional choices, optional conditions.
on_complete_flags	array	Flags set when the scene finishes, regardless of which choice was made.

2.3 Trigger Types

Trigger Type	Use Case
any_operative	Fires when ANY operative meets the threshold conditions. Most common for corruption milestones, stress breaks.
specific_operative	Fires for a named operative (by ID). Rare; used for story-critical characters.
relationship	Fires when the relationship score between two operatives crosses a threshold (positive > 60 for bonds, negative < -30 for tensions).
corporate_state	Fires based on board_confidence, faction hostility, directive_compliance. Used for corporate confrontation scenes.
post_mission	Fires immediately after a specific mission type completes. Used for debriefing, wet work aftermath.

2.4 Cast Resolution

Each cast slot resolves in order: (1) check the **role** tag against current operative fitness scores, (2) if no operative fits, use the **fallback** stat-based selector. Special values:

Value	Resolution
triggering_operative	The operative whose stats triggered the scene. Always resolves.
highest_social	Operative with the highest social stat among available cast.
lowest_loyalty	Operative with the lowest loyalty. Useful for tension scenes.
highest_conscience	Operative with the highest conscience. The moral center.
random_eligible	Random operative meeting minimum stat thresholds.
player_choice	Player selects from eligible operatives. Used sparingly for high-stakes scenes.

2.5 Effects System

Choices produce effects that modify game state. Effects are applied atomically when the player confirms a choice.

Effect Type	Syntax	Example
Stat delta	{ "target", "stat", "delta" }	conscience -8 on protagonist
Flag setting	"set_flags": ["flag_name"]	Gates future scene triggers
Corporate standing	"corporate_standing": { "delta": N }	Board confidence shift
Relationship	"relationship": { "target", "with", "delta" }	Bond/tension between two operatives

2.6 Gender Token System

CWC supports gendered text variation using inline tokens that the NarrativeDirector resolves at render time:

```
{gender:m|male text|female text}
```

The token checks the speaker's gender flag. Examples:

- {gender:m|He leaned back|She leaned back} in the chair.
- {gender:m|his|her} reputation preceded {gender:m|him|her}.
- {gender:m|Good man|Good woman}, the director said.

Tokens can appear anywhere in node text, choice labels, or effect descriptions. They cost nothing at runtime and are invisible to the player.

3. Use Case 1: Writing a Corruption Milestone Scene

3.1 Pick Your Threshold

Every corruption milestone scene is anchored to a specific moment in an operative's moral descent. The five corruption thresholds define the arc:

Corruption Index	Label	What It Means
20	Competent	The operative has proven effective. First kills behind them. Still processing.
40	Effective	Violence is becoming routine. Conscience is a voice, not a wall.
60	Feared	Reputation precedes the operative. Others avoid eye contact.
80	The Machine	The operative is an instrument. Feelings are operational noise.
95	Jenkins	The end state. The system made flesh. No way back.

3.2 Step-by-Step: The Third Wet Job Scene

Let's walk through building a complete corruption milestone scene from scratch.

Step 1: Define the Trigger

This scene fires when any operative completes their third assassination AND still has meaningful conscience remaining. The AND is crucial: an operative with zero conscience wouldn't experience this scene, because for them, there's nothing left to struggle with.

```
"trigger": {  
  "type": "any_operative",  
  "conditions": {  
    "wet_work_count": { "gte": 3 },  
    "conscience": { "gte": 40 }  
  }  
}
```

Step 2: Cast the Scene

Two roles: the protagonist (the operative who triggered) and an observer (the mirror role, someone perceptive enough to see what's happening).

```
"cast": {  
  "protagonist": "triggering_operative",  
  "observer": { "role": "mirror", "fallback": "highest_social" }  
}
```

Step 3: Write the Dialogue

The observer opens. They name what the protagonist can't. Keep it concrete, not philosophical. Cyberpunk register: technology is embodied, language is clipped, observation is clinical.

Writing Principle

The observer doesn't judge. They describe. 'Three names on your list now' is better than 'You've killed three people.' The first is a mirror; the second is a sermon.

Step 4: Design Choices

Every choice node needs 2-4 options. The design rule: **the cold option should be mechanically optimal but morally costly; the humane option should cost corporate standing but preserve the operative's soul.**

Choice	Mechanical Effect	Narrative Effect
'They were all necessary.'	conscience -8, stress -3 (relief through acceptance)	The operative stops fighting it. Easier now. That's the point.
'I need to talk to someone.'	stress -10, loyalty +3, corporate_standing -2	The operative reaches out. The corporation notices the hesitation.
'What's one more?'	conscience -12, stress -5, morale -4	Gallows humor. Cheapest coping mechanism. Costs the most long-term.

Step 5: Testing

Save the JSON file to **Data/Templates/corruption/**. In the debug console, use the state override to set any operative's wet_work_count to 3 and conscience to 45. Advance one cycle. The scene should trigger, cast resolve, and render. Verify all choices produce correct stat deltas.

4. Use Case 2: Writing an Operative Psychology Scene

Psychology scenes make the human simulation layer visible. They fire when an operative's internal state reaches a critical threshold, and they exist to let the player witness (and intervene in) an operative's psychological trajectory.

4.1 Trigger Types

Scene Type	Trigger Condition	Narrative Purpose
Stress Break	stress $\geq$ 80	The operative is cracking. Performance degrades. They may refuse orders, lash out, or shut down entirely.
Loyalty Test	loyalty $\leq$ 25	The operative is considering betrayal, defection, or going dark. A moment of decision.
Morale Collapse	morale $\leq$ 20	The operative has lost faith in the mission, the organization, or themselves. Different from stress: this is meaning, not pressure.
Conscience Crisis	conscience change $>$ 15 in one cycle	A sudden moral shock. Something happened that the operative can't rationalize away.

4.2 The Mirror Role

Psychology scenes almost always include a **mirror** character: someone who observes and names what the protagonist is experiencing. The mirror's function is narrative, not therapeutic. They don't fix anything. They make the internal external.

Good mirror dialogue:

- 'You haven't slept in three days. I can see it in how you're holding your sidearm.'
- 'The last time someone looked the way you look, they walked into the Baikonur wastes and didn't come back.'
- 'You're scaring the junior analysts. I don't think you know that.'

Bad mirror dialogue (avoid):

- 'Are you okay? You seem stressed.' (Too generic, too therapeutic)
- 'You need to talk to someone about your feelings.' (Anachronistic; wrong register)
- 'The mission is taking a toll on all of us.' (Deflection; loses the individual focus)

4.3 Voice by Archetype

An operative's background archetype should inflect how they express psychological distress. The same stress break plays differently depending on who's breaking.

Archetype	Voice Under Pressure
Street	Terse. Profane. Physical. Stress manifests as aggression, restlessness, picking fights. 'I didn't crawl out of the Reclamation Zones to die in a conference room.'
CorpoClimber	Controlled until it isn't. Stress manifests as over-performance, then sudden collapse. 'My quarterly review is in six days. I can hold it together for six days.'

Spook	Detached. Clinical. Stress manifests as hyper-analysis, paranoia, trust erosion. 'I've run the scenarios. Every exit has a watcher. Every watcher has a handler.'
Soldier	Stoic. Brief. Stress manifests as withdrawal, ritual behavior, cleaning weapons. 'Weapon's clean. Weapon's clean. Weapon's always clean.'
Fixer	Adaptive. Stress manifests as over-planning, contingency spirals, inability to commit. 'I need three more fallback positions before I can confirm the primary.'

5. Use Case 3: Writing a Mission Narrative Sequence

Mission sequences are the bridge between the strategy sandbox and the narrative layer. They play during active missions, giving the player narrative beats that modify mission parameters based on choices.

5.1 The 3-5 Node Structure

Every mission narrative follows a consistent structure:

Node	Function	Example
Briefing (1 node)	Set the scene. Establish stakes, target, location. The player should understand what they're walking into.	'The target is a data courier. Neural implant, encrypted payload. She'll be at the Kharkov exchange at 2200.'
Complication (1-2 nodes)	Something goes wrong or an unexpected element appears. Forces a choice that modifies mission parameters.	'There's a child in the room. The courier's daughter. She wasn't in the intel package.'
Resolution Modifier (1-2 nodes)	The player's choices here determine which skill weights apply to the final resolution roll.	'If you go loud, combat and demolitions. If you talk your way past, social and deception.'

5.2 The Approach System

Player choices during mission narratives shift which skills determine success. This is how narrative choices have mechanical teeth.

Approach	Skills Weighted
Stealth	infiltration, electronics, perception. Lowest stress gain on success. Highest on failure (blown cover is terrifying).
Social	deception, persuasion, social. Moderate stress. Requires the operative to maintain a persona.
Force	combat, demolitions, intimidation. Fastest resolution. Highest collateral. Guaranteed stress gain.

The approach isn't always a clean choice. The best mission narratives present situations where the optimal approach is ambiguous or where circumstances force a switch mid-mission.

5.3 Wet Work Decisions

The most important narrative moments in missions are **wet work decisions**: moments where the player must choose whether to kill. These are the bridge between strategy and narrative, the moments where `corruption_index` actually changes.

Design rules for wet work decisions:

- **Never make killing the only option.** There must always be an alternative, even if it's costly.
- **The alternative should cost something real.** More stress, lower mission success chance, corporate standing hit.

- **Name the target.** Even a one-line detail ('she has a data-port scar behind her left ear') makes the decision heavier.
- **Show the aftermath.** A post-mission node that acknowledges what happened. Not a lecture. A detail. Blood on a sleeve. A name on a list.

Design Principle

The player should never feel that killing was inevitable. They should feel that it was the easiest option. That distinction is the entire game.

6. Use Case 4: Writing Relationship and Corporate Scenes

6.1 Relationship Threshold Triggers

Relationships between operatives generate their own narrative moments. These scenes fire based on the accumulated relationship score between two specific operatives.

Threshold	Scene Type	Narrative Function
Bond > 60	Loyalty scene	Two operatives who trust each other. Creates vulnerability: one can be leveraged through the other.
Bond > 80	Partnership scene	Deep trust. They cover for each other. The player may have to choose between them.
Tension < -30	Confrontation scene	Open conflict. Affects team cohesion. May force reassignment or escalation.
Tension < -50	Breaking point	One will act against the other. Sabotage, betrayal, or a forcing function the player must resolve.

6.2 Corporate Confrontation Scenes

Corporate scenes trigger on organizational state, not individual stats. These are the moments when the boardroom becomes a battlefield.

Trigger	Scene
board_confidence < 30	The board is losing faith. A hearing. The player must justify their division's performance or face restructuring.
faction_hostility > 70	A rival faction moves against you. Corporate sabotage, resource raids, poached operatives.
directive_compliance < 40	You've been ignoring directives too long. An oversight committee appears. Consequences are coming.
board_confidence > 85	You're too successful. Other divisions notice. Jealousy manifests as political maneuvering.

6.3 Writing the Board Hearing

The board hearing is a signature CWC scene type. It inverts the power dynamic: the player, who commands operatives and authorizes wet work, is now answering to people who command them.

Structure:

7. Opening: A board member states the concern. Clinical corporate language. No emotion.
8. Evidence: Specific incidents cited. Missions that went sideways. Collateral. Budget overruns.
9. Player response: 2-4 choices. Defend aggressively (costs relationships, preserves standing). Accept blame (costs standing, preserves relationships). Deflect to a subordinate (costs loyalty, preserves everything else at someone else's expense).
10. Resolution: The board's verdict. Determined by accumulated corporate_standing, not just this scene's choice.

Writing Principle

The board hearing should feel like the inverse of a mission briefing. Same sterile language, same clinical detachment. The difference is that you're the target now.

7. The Corruption Arc: Design Philosophy

The corruption arc is CWC's central thesis: that systems corrupt the people who operate them, and that the corruption is invisible from the inside. The player should never receive a notification that says 'You are now corrupt.' They should realize it themselves, looking back.

7.1 The Five Milestone Thresholds

Index	Name	UI Change	Narrative Shift
20	Competent	Subtle color temperature shift. Slightly cooler palette.	Choices begin including a 'pragmatic' option that wasn't there before.
40	Effective	Further desaturation. Choice order begins shifting: 'cold' options appear first.	The observer role starts commenting on the protagonist's behavior.
60	Feared	Significant desaturation. Interface elements lose warmth.	NPCs react differently. Deference replaces collaboration. Fear replaces respect.
80	The Machine	Near-monochrome. Choices reorder: the ruthless option is always first.	The protagonist's internal voice changes. What was rationalization becomes statement of fact.
95	Jenkins	Full desaturation. The UI feels clinical, institutional.	The protagonist has become the system. No more internal conflict. Efficiency is its own justification.

7.2 The Design Principle: Show, Don't Tell

The game never labels the player's moral state. There is no 'Good/Evil' meter, no alignment notification, no achievement for reaching a corruption threshold. Instead:

- **Choices reorder.** At corruption 0, the humane choice is first. By corruption 80, the ruthless choice is first. The player notices they're clicking the top option more often. They realize the top option has changed.
- **Language shifts.** Early-game choice labels say 'Eliminate the target.' Late-game labels say 'Clean up.' The euphemism treadmill works on the player just as it works on the operative.
- **The world reacts.** Junior operatives flinch. Allies stop making small talk. The canteen goes quiet when you enter. These are atmosphere nodes, not plot nodes.
- **The UI desaturates.** Color drains from the interface at each threshold. It happens slowly enough that the player adapts. That adaptation IS the point.

7.3 Writing Across the Arc

When writing scenes that can trigger at multiple corruption levels, use conditional text or separate scene variants. A stress break at corruption 20 looks completely different from a stress break at corruption 80:

Corruption Level	Same Scene, Different Expression
20 (Competent)	'I can't sleep. Every time I close my eyes I see his face. The way he looked at me when he knew.'

60 (Feared)	'I can't sleep. Operational review is at 0800 and the after-action report has inconsistencies I need to resolve before the committee sees it.'
80 (The Machine)	'I can't sleep. Suboptimal. I'll requisition pharmacological intervention and be operational by 0600.'

Same trigger (stress $\geq$ 80). Same surface behavior (insomnia). Completely different interiority. The first operative is haunted. The third has replaced their humanity with process.

8. Cyberpunk Register Guide

CWC's dialogue should feel like it exists in a world where neural implants are as mundane as smartphones, where corporations have replaced nation-states as the primary organizing force, and where language has evolved to reflect both of those realities.

8.1 Technology Is Embodied

In CWC's world, technology is not external. It's installed. Characters don't 'use' a computer; they have cortical stacks. They don't 'check their phone'; they parse their feed overlay. Write technology as body, not tool.

Instead Of	Write
'She checked her phone'	'Her optic overlay flickered with incoming traffic'
'He accessed the database'	'He dropped into the datawell, cortex-first'
'The security system activated'	'The building's immune response kicked in'
'She sent a message'	'She squirted a compressed burst through the mesh'
'He turned off his implant'	'He went dark. Like unplugging a limb.'

8.2 Corporate Euphemism

The corporation never says what it means. This is both realistic and thematically essential: the euphemism treadmill is how the system hides its own violence from itself.

Real Meaning	Corporate Language
Assassination	Asset reallocation / Permanent contract termination
Firing an employee	Workforce optimization event
Covering up a murder	Narrative management / Incident containment
Torture / interrogation	Enhanced information extraction
Civilian casualties	Collateral inefficiency
Blackmail	Leverage-based negotiation
Betrayal	Strategic realignment of loyalties

8.3 Street-Level Contrast

Operatives from street backgrounds don't use corporate euphemism. They call killing 'killing.' This contrast is deliberately designed: it makes the corporate language feel more alien, more sanitized, more complicit.

Street register: short sentences, concrete nouns, physical verbs. 'I dropped him in the alley behind the noodle place. He bled out before his immune system could close the wound.'

Corporate register: complex clauses, abstract nouns, passive voice. 'The asset was neutralized in a low-visibility environment. Biological termination occurred prior to automated medical intervention.'

8.4 Architecture of Power

Describe spaces by access level, not aesthetics. A beautiful office is less important than the fact that it requires Level 4 clearance to enter. Security tiers, access codes, and surveillance coverage tell the player more about power dynamics than any description of marble floors.

8.5 Information as Currency

In CWC's world, data has weight. Blackmail files are assets. Access codes are leverage. A conversation isn't just dialogue; it's a negotiation where information is the commodity. Write scenes where characters trade knowledge, withhold it strategically, or weaponize it.

9. Gender Expression Guide

CWC supports both male and female player characters. The gender system is designed to add texture, not double the content.

9.1 Token Syntax

```
{gender:m|male text|female text}
```

The token checks the **m** flag (male = true). If true, renders the first option. If false, renders the second. Tokens can be nested in any text field.

9.2 Where Gender Matters

Gender expression should inflect scenes at 15-20 key moments across the full game, not in every line of dialogue. The moments where it matters are moments of **power dynamics**: how authority is asserted, how it's perceived, how resistance manifests.

- **Authority assertion:** A male operative commands; a female operative commands and is questioned about it. Not because the game is making a statement, but because the world is.
- **Power dynamics:** Corporate hierarchy perceives gender. A male director is 'decisive'; a female director is 'aggressive.' The same behavior, different reception.
- **Physical descriptions:** Sparing, but gendered where it adds character. Posture, movement, how space is occupied.

9.3 The Two Protagonist Modes

CWC's protagonist has two expressive modes, loosely modeled on archetypes from prestige television:

Mode	Description
Draper Energy (Male Default)	Command through mastery. The room orients around the protagonist because competence is assumed. Power is quiet, gravitational. The danger is complacency: the world bends, and bending the world becomes invisible.
Stout Energy (Female Default)	Weaponized competence. The protagonist has armor made of being better than everyone in the room, and she knows it. The system wasn't built for her, so she became the system. The danger is the same: becoming the thing you learned to beat.

Both modes reach the same destination: the corruption arc doesn't change based on gender. What changes is the texture of the journey. The system corrupts differently, but it corrupts completely.

9.4 Don't Double Content

The gender system is designed for texture, not bifurcation. Do not write separate scenes for male and female characters. Write one scene with gender tokens at the moments where the difference creates meaning. Most dialogue is gender-neutral. Most choices are gender-neutral. The tokens exist for the 15-20 moments where the gendered experience of power adds something the gender-neutral version can't.

10. Quick Reference: Scene Template Checklist

10.1 Blank Scene Template

```
{
  "id": "",
  "priority": 50,
  "cooldown_cycles": 0,
  "trigger": {
    "type": "",
    "conditions": {}
  },
  "cast": {
    "protagonist": "triggering_operative"
  },
  "nodes": [
    {
      "id": "opening",
      "speaker": "protagonist",
      "portrait_mood": "neutral",
      "text": "",
      "choices": [
        {
          "label": "",
          "effects": [],
          "set_flags": [],
          "next": ""
        }
      ]
    }
  ],
  "on_complete_flags": []
}
```

10.2 Pre-Submission Checklist

Check	Question	If No
Trigger	Is the trigger condition specific enough to fire at the right moment?	Add additional stat conditions or flags to narrow the trigger window.
Cast Fallbacks	Does every cast role have a fallback selector?	Add fallback: highest_social, lowest_loyalty, etc.

Choice Consequences	Do all choices produce meaningfully different effects?	If two choices lead to the same outcome, merge them or differentiate the effects.
Cost	Does at least one choice cost the player something real?	Add a corporate_standing hit, a loyalty cost, or a stress gain.
Observer Witness	Is the observer/mirror flagged if present?	Set the 'witnessed_by_mirror' flag so future scenes know someone saw.
Gender Tokens	Are gender tokens placed at power-dynamic moments (if applicable)?	Add tokens where authority assertion, perception, or resistance differs.
Flags	Does on_complete_flags prevent re-triggering?	Add a unique flag to on_complete_flags.
Cooldown	Is cooldown_cycles set appropriately?	0 for one-shot scenes; 3-5 for repeatable scenes.
Priority	Is priority set relative to other scenes in this trigger range?	Higher priority (80-100) for milestone scenes; lower (30-50) for atmosphere.
Voice	Does the dialogue match the operative archetype's register?	Review Section 4.3 and Section 8 for register guidance.

10.3 File Naming Convention

Pattern	Example
corruption_{milestone}_{descriptor}	corruption_third_wet_job.json
psych_{type}_{descriptor}	psych_stress_break_spook.json
mission_{mission_type}_{phase}	mission_extraction_complication.json
relationship_{bond tension}_{descriptor}	relationship_bond_partnership.json
corporate_{event_type}	corporate_board_hearing.json

10.4 Directory Structure

```
Data/Templates/
  corruption/      -- Corruption milestone scenes
  psychology/     -- Stress, loyalty, morale scenes
  missions/       -- Mission narrative sequences
  relationships/   -- Bond and tension scenes
  corporate/      -- Board, faction, directive scenes
  atmosphere/     -- Non-triggered ambient scenes
```

Save each scene as a single JSON file in the appropriate subdirectory. The NarrativeDirector scans all subdirectories recursively on startup and hot-reloads on file change during debug mode.